

# MD AMIR KHAN

Applied AI Engineer · LLM Systems & Quantitative Analytics

[mkhan37@stevens.edu](mailto:mkhan37@stevens.edu) · Stamford, CT · [github.com/Mkhan2317](https://github.com/Mkhan2317) · [linkedin.com/in/amirkhan2317](https://linkedin.com/in/amirkhan2317) · [mkhan2317.github.io](https://mkhan2317.github.io)

## SUMMARY

Applied AI engineer building production **LLM systems** — structured extraction, agentic RAG, multi-model inference — and the **ML forecasting stack** that consumes them, delivered end-to-end from FastAPI back-ends to Next.js / React front-ends and running daily on the DXT Commodities gas & power trading desk.

## EDUCATION

### Stevens Institute of Technology

Hoboken, NJ, USA

Master of Science in Financial Engineering & Analytics (STEM-designated)

### North South University

Dhaka, Bangladesh

Bachelor of Business Administration · Major: Finance · Minor: Mathematics

## EXPERIENCE

### DXT Commodities North America

Stamford, CT

Applied AI Engineer & Quantitative Analyst — Market Fundamentals (LNG & Power)

March 2026 – Present

- **Pipeline transparency platform — LLM structured extraction at production scale.** Built the scraping and LLM extraction pipeline reading maintenance and capacity notices from 19 U.S. interstate gas pipelines across 8 different portal architectures (15+ format-specific scrapers: HTML, PDF, Excel, ASP.NET with TLS impersonation for WAF-protected sites). **Structured extraction on the Claude API** turns free-form notices into typed capacity-impact records with a caching layer that cuts LLM calls on unchanged input. FastAPI microservices + SQL Server; REST APIs; tracks 40+ active and 198+ upcoming restrictions on a peak day.
- **WattsApp — full-stack PJM analytics deployed daily to the power desk.** LLM-powered natural-language query view, plus overview, thermal, non-thermal, plants, and insights modules covering **4,305 units / 232.5 GW nameplate** across 2,262 plants in 13 states + DC. Next.js + React + TypeScript, FastAPI, SQL Server, Docker, REST APIs.
- **U.S. LNG feed-gas forecasting.** End-to-end forecasting pipeline covering **126 export trains at 13 terminals (~22,600 MMcf/d nameplate)**. Three-model validation framework; **under 3% MAPE** on an 8-month out-of-sample holdout.
- **U.S. natural-gas demand forecasts.** 48 models — 12 regions × 4 end-use sectors — producing daily 30-day forecasts for the Lower 48. XGBoost for weather-sensitive sectors, Ridge for the slower ones; walk-forward cross-validation. Driven by NWS temperature forecasts and Fed industrial-production data, trained on EIA state-level consumption.
- **Wattson — multi-ISO power-price analytics.** Live platform for PJM, ERCOT, NYISO, MISO, and ISO-NE: historical LMPs by energy / congestion / marginal-loss component, forward-contract settlements with strip aggregation, and a Price ⇌ Heat-Rate toggle connecting power and gas markets.
- **Real-time gas-production model & Permian Basin intelligence.** OLS/Ridge scaling framework combining licensed pipeline-nominations data with EIA-914 to produce a current-month U.S. production estimate inside the two-month EIA reporting lag. Permian codebase covers daily production, egress across the four major pipelines, and Waha basis pricing.
- **Pipeline Force Majeure alerts & PEPCO nodal-basis screen.** Real-time alerting service polls electronic bulletin boards across **32 U.S. interstate gas pipelines** every 5 minutes and pushes notices to Microsoft Teams (response time: hours → minutes). PEPCO FTR-bidding screen covers **~8.77 M hourly observations** across 61 nodes and 8+ years of PJM day-ahead settlement data.

### Stevens Institute of Technology

Hoboken, NJ

Research Assistant — LLM-Generated Views in a Black-Litterman Portfolio (MSc Thesis)

January 2024 – December 2025

- Designed and built an **LLM-powered forecasting system** treating language models as queryable, uncertainty-aware forecasters rather than static tools — a case study in structured output extraction, multi-model evaluation, and confidence-aware system design.
- Built a **multi-LLM inference pipeline** (GPT-4o-mini, LLaMA-3.1, Gemma, Qwen-2) using **vLLM** for high-throughput serving; structured prompts convert price, sector, and news data into scalar return forecasts.
- Solved the **uncertainty-quantification problem** by repeated querying and treating output dispersion as a Bayesian confidence signal — sample mean fed as the view, sample variance as view uncertainty, into a downstream **Black-Litterman** update rather than being discarded.
- Full backtesting pipeline: 50 largest S&P 500 names, **Feb 2024 – Jan 2025**, biweekly rebalancing, benchmarked against 3 baseline strategies; **best model (LLaMA-3.1) beat the S&P 500 benchmark — Sharpe 1.29 vs. 0.87.**

### Stevens Institute of Technology

Hoboken, NJ

Quantitative Research Assistant — School of Business

April 2025 – February 2026

- Implemented the  $\varepsilon$ -subdivision **Robust PCA framework for dynamic factor portfolios** in Python with Prof. Papa Momar

- Ndiaye — block decomposition of the eigenspectrum against a tolerance  $\varepsilon$ , Gram–Schmidt construction of the closest orthonormal basis to the prior period’s factors, and rupture-detection that resets factor tracking when block-mean eigenvalues shift beyond a threshold  $\delta$ . Added K-Means clustering on eigenvector features as an independent signal of factor-structure change.
- Validated on  $\sim 6.5$  years of daily returns across the 11 S&P 500 GICS sector ETFs, spanning the COVID-19 shock and the post-pandemic inflation cycle. Fed robust and standard covariance estimates into a mean-variance optimizer; the robust approach stabilized factors at portfolio volatility essentially identical to standard PCA.

## TEACHING EXPERIENCE

---

### Stevens Institute of Technology

*Hoboken, NJ*

*Teaching Assistant — Quantitative Portfolio Theory (Financial Engineering program)*

*September 2025 – December 2025*

## TECHNICAL SKILLS

---

|                                |                                                                                                                                                                                                                                                                |
|--------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>Full-Stack &amp; Infra:</b> | FastAPI · Next.js · React · TypeScript · SQLAlchemy · SQL Server · Docker · AWS · Microservices · REST APIs · Git / GitHub / GitLab                                                                                                                            |
| <b>AI / LLM Engineering:</b>   | LangChain · LangGraph · LangSmith · RAG (HyDE, CRAG, Self-RAG, Graph RAG, Agentic RAG) · LLM Agents & Subagents · Structured Output Extraction · Prompt Engineering · Multi-LLM Serving (vLLM) · Anthropic Claude & OpenAI APIs · Model Context Protocol (MCP) |
| <b>Languages:</b>              | Python · SQL · TypeScript · JavaScript                                                                                                                                                                                                                         |
| <b>Machine Learning:</b>       | Scikit-learn · XGBoost · Ridge / OLS · Statsmodels · Pandas · NumPy · Time-Series Forecasting · Walk-Forward Cross-Validation · Robust PCA · Black–Litterman                                                                                                   |
| <b>Domain (Energy):</b>        | U.S. Natural Gas & Power Fundamentals · LNG Feed-Gas · Pipeline Capacity Analysis · ISOs / RTOs (PJM, ERCOT, NYISO, MISO, ISO-NE) · Wholesale LMPs · FTR Markets · Heat-Rate Analysis                                                                          |

## CERTIFICATIONS & AWARDS

---

- Anthropic Education — Certified Track** (Anthropic, 2026): Claude API, Claude Code, Model Context Protocol (Intro + Advanced), Agent Skills, Subagents.
- Advanced RAG** (CampusX, 2026) · **AI Engineer Bootcamp 2026: LLMs, RAG, Agents, Vector DBs** (Udemy, 2026) · **Vector Databases: Fundamentals to Production** (Udemy, 2026).
- Full-Stack Web Track** (Modern JavaScript, TypeScript, React, Next.js, Node/Microservices, Full-Stack Bootcamp; Udemy, 2026) · **Developer Foundations** (Linux CLI, Git & GitHub; Udemy, 2026).
- Probability — The Science of Uncertainty and Data** (MIT / edX, 2022) · **Python and Statistics for Financial Analysis** (Coursera, 2022) · **Akuna Capital Options 101** (2025) · **Algorithmic Trading with Python, ChatGPT, ML** (Udemy, 2025) · **FastAPI Complete Course** (Udemy, 2025).
- Student Member, **CFA Society New York**.